



Production, Manufacturing and Logistics

New benchmark instances for the Capacitated Vehicle Routing Problem



Eduardo Uchoa^a, Diego Pecin^b, Artur Pessoa^a, Marcus Poggi^b, Thibaut Vidal^b,
Anand Subramanian^{c,*}

^a Departamento de Engenharia de Produção, Universidade Federal Fluminense, Rua Passo da Pátria, 156 São Domingos, Bloco E – 4o andar, Niterói – RJ, 24210-240, Brazil

^b Departamento de Informática, Pontifícia Universidade Católica do Rio de Janeiro, Rua Marquês de São Vicente, 225 Gávea, Rio de Janeiro – RJ, 22451-900, Brazil

^c Departamento de Sistemas de Computação Centro de Informática, Universidade Federal da Paraíba, Rua dos Escoteiros, Mangabeira, João Pessoa – PB, 58058-600, Brazil

ARTICLE INFO

Article history:

Received 11 March 2016

Accepted 4 August 2016

Available online 10 August 2016

Keywords:

Routing

Benchmark instances

Experimental analysis of algorithms

ABSTRACT

The recent research on the CVRP is being slowed down by the lack of a good set of benchmark instances. The existing sets suffer from at least one of the following drawbacks: (i) became too easy for current algorithms; (ii) are too artificial; (iii) are too homogeneous, not covering the wide range of characteristics found in real applications. We propose a new set of 100 instances ranging from 100 to 1000 customers, designed in order to provide a more comprehensive and balanced experimental setting. Moreover, the same generating scheme was also used to provide an extended benchmark of 600 instances. In addition to having a greater discriminating ability to identify “which algorithm is better”, these new benchmarks should also allow for a deeper statistical analysis of the performance of an algorithm. In particular, they will enable one to investigate how the characteristics of an instance affect its performance. We report such an analysis on state-of-the-art exact and heuristic methods.

© 2016 Elsevier B.V. All rights reserved.

1. Introduction

The Vehicle Routing Problem (VRP) is among the most widely studied problems in the fields of operations research and combinatorial optimization. Its relevance stems from its direct application in the real world systems that distribute goods and provide services, vital to the modern economies. Reflecting the large variety of conditions present in those systems, the VRP literature is spread into dozens of variants. For example, there are variants that consider time windows, multiple depots, mixed vehicle fleet, split deliveries, pickups and deliveries, precedences, complex loading constraints, etc.

In the vast landscape of VRP variants, the Capacitated VRP (CVRP) occupies a central position. It was defined by Dantzig and Ramser (1959) as follows. The input consists of a set of $n + 1$

points, a depot and n customers; an $(n + 1) \times (n + 1)$ matrix $d = [d_{ij}]$ with the distances between every pair of points i and j ; an n -dimensional demand vector $q = [q_i]$ giving the amount to be delivered to customer i ; and a vehicle capacity Q . A solution is a set of routes, starting and ending at the depot, that visit every customer exactly once in such a way that the sum of the demands of the customers of each route does not exceed the vehicle capacity. The objective is to find a solution with minimum total route distance. Some authors also assume that the number of routes is fixed to an additional input number K (almost always defined as the minimum possible number of routes, K_{min}).

The CVRP plays a particular role on VRP algorithmic research, for both exact and heuristics methods. Being the most basic variant, it is a natural testbed for trying new ideas. Its relative simplicity allows cleaner descriptions and implementations, without the additional conceptual burden necessary to handle more complex variants. Successful ideas for the CVRP are often later extended to more complex variants. For example, the classical CVRP heuristic by Clarke and Wright (1964) was adapted for the VRP with Time Windows (Solomon, 1987) and for many other variants, as surveyed in Rand (2009).

* Corresponding author.

E-mail addresses: uchoa@producao.uff.br (E. Uchoa), dpecin@inf.puc-rio.br (D. Pecin), artur@producao.uff.br (A. Pessoa), poggi@inf.puc-rio.br (M. Poggi), vidalt@inf.puc-rio.br (T. Vidal), anand@ci.ufpb.br, anandsubraman@gmail.com (A. Subramanian).

The present work is motivated by the claim that the recent research on the CVRP is being hampered by the lack of a well-established set of benchmark instances able to push the limits of the state-of-the-art algorithms. We examine this claim in the following.

Since Augerat et al. (1995), almost all new exact methods for the CVRP have been tested over instances from six different classes. Classes A, B and P were proposed in Augerat et al. (1995), while classes E, F and M were proposed by Christofides and Eilon (1969), by Fisher (1994) and by Christofides, Mingozzi, and Toth (1979), respectively. They are often referred to as the ABEFMP instances. Nearly all those instances are Euclidean. In the literature of exact methods it became usual to follow the TSPLIB convention of rounding distances to the nearest integer (Reinelt, 1991) and also to fix the number of routes.

The ABEFMP instances provided a good benchmark until a decade ago. Although 90 out of its 95 instances have no more than 100 customers, several such instances were still very challenging for the branch-and-cut algorithms (like Achuthan, Caccetta, & Hill, 2003; Araque, Kudva, Morin, & Pekny, 1994; Augerat et al., 1995; Blasum & Hochstättler, 2000; Lysgaard, Letchford, & Eglese, 2004; Ralphs, Kopman, Pulleyblank, & Trotter, 2003; Wenger, 2003), which prevailed at that time. Then, the Branch-Cut-and-Price of Fukasawa et al. (2006) solved all its instances with up to 100 customers, as well as instances M-n121-k7 and F-n135-k7 (the latter was already solved since Augerat et al., 1995). Only instances M-n151-k12, M-n200-k16 and M-n200-k17 remained unsolved. From that moment, the benchmark was not satisfactory anymore. While the majority of its instances became quite easy, the now more interesting range of 101–200 customers was very thinly populated. In particular, the column and cut generation algorithms of Baldacci, Christofides, and Mingozzi (2008) and Baldacci, Mingozzi, and Roberti (2011) significantly reduced the CPU time required to solve many ABEFMP instances, but could not solve the three larger M instances. Later, the algorithms of Contardo and Martinelli (2014) and Røpke (2012) were capable of solving M-n151-k12. Very recently, the algorithm of Pecin, Pessoa, Poggi, and Uchoa (2014) solved the last two open instances in the benchmark.

Since 1980's, almost all papers proposing new heuristic and metaheuristic methods for the CVRP reported results on a subset of the instances by Christofides and Eilon (1969) and Christofides et al. (1979). In this literature, it is usual to follow the convention of not rounding the Euclidean distances and not to fix the number of routes. This classical benchmark is also exhausted, since most recent heuristics systematically find the best known solutions on nearly all instances. In fact, Rochat and Taillard (1995) already published in 1995 what we now know to be the optimal solutions, except for a single instance with 199 customers, where the reported solution was only 0.012% off-optimal.

The more recent set of instances by Golden, Wasil, Kelly, and Chao (1998) is now the second most frequently used benchmark. Having larger instances, ranging from 240 to 483 customers, it still has a good discriminating power. In fact, the heuristics of Nagata and Bräysy (2009) and Vidal, Crainic, Gendreau, Lahrichi, and Rei (2012) are considered the best available for the CVRP basically due to their superior performance on Golden's instances. Nevertheless, we believe that those instances are not sufficient for a good benchmark. A first drawback is their artificiality. In all instances the customers are positioned in concentric geometric figures, either circles, squares or six-pointed stars. The demands also follow very symmetric patterns. As a result, the solution space is partitioned into groups of equivalent solutions obtainable by rotations and flippings around several axis of symmetry. The second drawback is their relative homogeneity. For example, there are no instances with clusters of customers.

This work proposes a main set of 100 instances, ranging from 100 to 1000 customers, intended to be used by both exact and heuristic methods in the next years, thus promoting a cross-fertilization between both such methods. These instances improve upon the existing benchmarks because they are less artificial, avoiding symmetric patterns that are unlikely to appear in real data, and less homogeneous, since they vary not only in terms of size but also according to four parameters which independently affect the depot position, the customer positions, the demands, and the average route length. This work also presents new analyses of recent state-of-the-art heuristic and exact methods. The new benchmark allows to better assess the performance of solution approaches in relation to different instance features, makes it possible to better discriminate competing algorithms, and provides new insights on the benefits and drawbacks of such methods.

The number of instances was chosen to be not so large and still provide a more comprehensive and balanced experimental setting. For those interested in performing a deeper statistical analysis of the performance of an algorithm (at the expense of a substantially larger CPU time effort), we also propose an extended set containing 600 instances so as to enable one to closely investigate how the characteristics of an instance affect its performance. We remark that there are classical papers where the main contribution is proposing instances, like Reinelt (1991) that describes the TSPLIB benchmark still in use for the traveling salesman problem. However, recent papers with that spirit (like Vallada, Ruiz, & Framinan, 2015) also present extensive computational experiments with state-of-the-art methods, in order to prove that the benchmark indeed has the desired discriminating power and also provide good lower and upper bounds for each instance.

The remainder of the paper is organized as follows. Section 2 explains the procedure developed for generating new CVRP instances. Section 3 presents the main benchmark set and include results from state-of-the-art heuristic and exact methods on each of those 100 instances. Section 4 presents the extended benchmark and includes a more sophisticated statistical analysis of such methods over those instances. Section 5 briefly describes the features of the new CVRPLIB web site where all instances mentioned in this paper are available. Section 6 presents the concluding remarks of this work. Finally, an appendix presents detailed information on all existing CVRP benchmark instances, explaining their origin and providing their optimal/best known solution values.

2. Generation mechanism

This section describes how the instances in the proposed CVRP benchmarks were generated. As happens in almost all the existing instances, the distances are two-dimensional Euclidean. Depot and costumers have integer coordinates corresponding to points in a $[0, 1000] \times [0, 1000]$ grid. Each instance is characterized by the following attributes: number of customers, depot positioning, customer positioning, demand distribution, and average route size. The possible values of each attribute and their effect in the generation are described in the next section.

2.1. Instance attributes

2.1.1. Depot positioning

Three different positions for the depot are considered:

Central (C) – depot in the center of the grid, point (500,500).

Eccentric (E) – depot in the corner of the grid, point (0,0).

Random (R) – depot in a random point of the grid.

The TC, TE and TR instances (Gouveia, 1996), used as benchmark on rooted network design problems, present similar alternatives for root positioning.

2.1.2. Customer positioning

Three alternatives for customer positioning are considered, following the R, C and RC instance classes of the Solomon set for the VRPTW (Solomon, 1987).

Random (R) – all customers are positioned in random points of the grid.

Clustered (C) – at first, a number S of customers that will act as cluster seeds is picked from a uniform discrete distribution $UD[3,8]$. Next, the S seeds are randomly positioned in the grid. The seeds will then attract, with an exponential decay, the other $n - S$ customers: the probability for a point p in the grid to receive a customer is proportional to

$$\sum_{s=1}^S \exp(-d(p, s)/40),$$

where $d(p, s)$ is the distance between p and seed s . The divisor 40 in the above formula was chosen after a number of experiments. Smaller values lead to excessively dense and isolated clusters. On the other hand, larger divisors result in clusters that are too sparse and mingled. When two or more seeds happen to be close, their combined attraction is likely to form a single larger cluster around them. This means that the size of the clusters may differ significantly. When two seeds are a little more apart, their clusters will not coalesce but a “bridge” of customers between them may appear. A seed that is sufficiently apart from other seeds will form an isolated cluster. The overall clustering scheme was devised to mimic the densities found in some large urban agglomerations that have grown from more or less isolated original nucleus (the seeds).

Random-clustered (RC) – half of the customers are clustered by the above described scheme, the remaining customers are randomly positioned.

It should be noted that superpositions are not allowed, all customers and the depot are located in distinct points on the grid.

2.1.3. Demand distribution

Seven options of demand distributions have been selected in these instances, with either specific values, or randomly generated with a different range and/or coefficient of variation (CV), defined as the ratio between the standard deviation and the mean:

Unitary (U) – all demands have value 1.

Small values, large CV (1–10) – demands from $UD[1,10]$.

Small values, small CV (5–10) – demands from $UD[5,10]$.

Large values, large CV (1–100) – demands from $UD[1,100]$.

Large values, small CV (50–100) – demands from $UD[50,100]$.

Depending on quadrant (Q) – demands taken from $UD[1,50]$ if a customer is in an even quadrant (with respect to point (500,500)), and from $UD[51,100]$ otherwise. This kind of demand distribution leads to solutions containing some routes that are significantly longer, in terms of number of served customers, than others.

Many small values, few large values (SL) – most demands (70–95% of the customers) are taken from $UD[1,10]$, the remaining demands are taken from $UD[50,100]$.

2.1.4. Average route size

The previous experience of the authors with CVRP algorithms indicated that the value n/K_{min} , the average route size (assuming solutions with the minimum possible number of routes) has a large impact on the performance of current exact methods. The impact of this attribute on heuristic methods is not so pronounced, but is still quite significant. In the case of exact algorithms, modern Branch-Cut-and-Price algorithms perform much better than

branch-and-cut algorithms on almost all instances (see Poggi and Uchoa, 2014 for a detailed discussion). However, there are a few instances with long routes (like F-n135-k7, see Pecin et al., 2014) where a branch-and-cut similar to the one presented in Lysgaard et al. (2004), running in a modern machine, still obtain better times.

As a matter of fact, we cannot make a general statement that instances with shorter routes are easier than instances with longer routes or the opposite. What happens is that some methods are more suited for short routes and other methods for longer routes. Therefore, a comprehensive benchmark set should contain instances where this attribute vary over a wide range of values. Yet, generating an instance with exactly a given value of n/K_{min} would be difficult, since even computing K_{min} requires the solution of a bin-packing problem. The generator actually uses as attribute a value r representing the desired value of n/K_{min} . This causes the instance capacity to be defined as:

$$Q = \left\lceil \frac{r \sum_{i=1}^n q_i}{n} \right\rceil,$$

where the q vector represents the demands, already obtained according to the specified demand distribution. As $\sum_{i=1}^n q_i/Q$ is usually a good lower bound on K_{min} , instances have values of n/K_{min} that are sufficiently close to r .

2.2. Two decisions on conventions

2.2.1. Rounding the distances or not?

The Euclidean distances are rounded to the nearest integer in the literature on exact methods, following the TSPLIB convention (Reinelt, 1991). On the other hand, the distances are seldom rounded in the literature on heuristics.

Advantages of rounding– most mathematical programming based algorithms (including standard MIP solvers) have a limited optimality precision. For example, CPLEX 12.5 default precision is only 10^{-4} (0.01%). It is possible to increase the precision up to a point by adjusting parameters (say, 10^{-6} or 10^{-7}). Going further requires special software, using more bits in the floating-point numbers or even exact rational arithmetic (Cook, Koch, Steffy, & Wolter, 2011), that is not easily available or implementable. The practice of distance rounding is convenient for avoiding those pitfalls and usually makes the optimal values found by exact methods based on standard mathematical programming software reliable. Remark that the practice of not rounding, but only reporting two decimal places does not solve the problem. For example, an algorithm with a precision of 10^{-6} may declare a solution having value 853.2351 (published as 853.24) as optimal. Later, someone finds a solution with value 853.2349 and publishes 853.23.

Disadvantages of rounding– a benchmark of rounded instances has less power for comparing competing algorithms. Especially for heuristics, the search space is formed by a relatively small number of plateaus, i.e., sets of solutions with the same value. Guiding the search on such plateaus is usually done by trial and error since there is no indication that the distance is – even slightly – reduced. In practice, there may be several distinct optimal solutions, leading to more frequent ties between competitors. This effect is quite significant on ABEFMP instances, where the optimal solution values have magnitudes around 10^3 . In addition, some people claim (but never publish, this is part of the community folklore) that rounding can artificially enhance the performance of exact methods. For example, if an upper bound of 1000 is known, a branch-and-bound node with lower bound $999 + \epsilon$ can be fathomed. This effect is significant on ABEFMP instances, enough to make some algorithms to run at least twice as fast.

We took the decision that the newly proposed benchmark set will follow the TSPLIB convention of rounding distances. Nevertheless, the

instances were devised in order to minimize the above mentioned disadvantages. The use of a $[0, 1000] \times [0, 1000]$ grid (instead of the $[0, 100] \times [0, 100]$ grids of most ABEFMP instances) makes the optimal solutions to have magnitudes between 10^4 and 10^5 . This is still safe for exact methods having a precision of 10^{-6} . However, the plateau effect and the artificial enhancement of exact methods are much reduced.

2.2.2. Fix the number of routes or not?

The literature on exact methods usually follows the convention of fixing the number of routes to a value K . Except for instance M-n200-k17, K is always set to K_{min} . The standard explanation for that fixing is that K represents the number of vehicles available at the depot. A more sophisticated explanation is that fixing the number of routes to the minimum is an indirect way of minimizing the fixed costs for using a vehicle. We do not think that this is necessarily true: since the CVRP definition allows solutions containing routes that are much shorter than others, why not assigning the two shortest routes to the same vehicle? In fact, the CVRP may be used to model a real-world situation where a single vehicle will perform all routes in sequence. In other words, CVRP routes do not need to correspond to vehicles. Based on that reasoning, we decided that the newly proposed benchmark set will follow the convention of not fixing the number of routes. Therefore, the number K_{min} indicated in each instance should be taken only as a lower bound on the number of routes in a solution. A second reason for taking that decision is that the number of routes was not fixed in the original CVRP definition (Dantzig & Ramser, 1959).

3. Main benchmark

A benchmark set with instances corresponding to the Cartesian product of so many parameter values (even restricting the “continuous” parameters n and r to a reasonably small number of values) would be huge. Thus, we generated a sample of 100 instances, presented in Tables 1–3. We believe that this number of instances is large enough to obtain the desired level of diversification, but still small enough to allow future users of the benchmark to report detailed results for each instance. We now explain how the attribute values of those 100 instances were obtained.

The instances are ordered by the number of customers n . For the first 50 instances, ranging from 100 to 330 customers, n is increased in linear steps. For the last 50 instances, n is increased in exponential steps from 335 to 1000. For the time being, instances with around 200 customers can already be hard for exact methods, and larger instances are sufficiently challenging to highlight significant quality differences between competing heuristics.

The set of values of r were taken from a continuous triangular distribution $T(3,6,25)$ (minimum 3, mode 6 and maximum 25). The 100 values of r were partitioned into quintiles, corresponding to very small routes, small routes, medium routes, long routes and very long routes. The k th instance receives an r from the $((k-1) \bmod 5) + 1$ quintile. In this way, one guarantees that every set containing $5t$ instances with consecutive values of n will have exactly t instances from each quintile. It can be observed at the end of Table 3 that the resulting set of n/K_{min} values have minimum 3.0, maximum 24.4, median 9.8 and average 11.1.

A random permutation of the 3 possible values for the depot positioning attribute (C, E and R) provides the values for the first 3 instances. Another random permutation gives the values for the next 3 instances and so on. A similar scheme is used for obtaining the values for the customer positioning and demand distribution attributes. This ensures that every subset of the instances with consecutive values of n will have a near-balanced number of instances with the same value of an attribute.

The name of an instance follows the ABEFMP standard, and has a format X-nA-kB, where A represents $n+1$, the number of points in the instance including the depot, and B is the minimum possible number of routes K_{min} , calculated by solving a bin-packing problem.

In order to provide an initial set of results to the new benchmark, experiments have been conducted with recent state-of-the-art heuristic and exact methods. The experiments also provide an assessment of the level of difficulty of the proposed instances and their capacity of discriminating competing algorithms. Those results are reported jointly in Tables 1–3, which describe the characteristics of each instance, the average solution quality (Avg), best solutions quality (Best) and average CPU time ($T(\text{minutes})$) of two state-of-the-art metaheuristics, the root lower bound (RLB), root CPU time ($RT(\text{minutes})$), number of search nodes (Nds) and overall time ($T(\text{minutes})$) of a recent state-of-the-art exact method, and finally the cost (Value) and number of vehicles (NV) of the best solution (BKS) ever found since the instances were created, including preliminary runs of those heuristics with alternative parameterizations. Table 3 also reports additional statistics on instances characteristics and results, such as the minimum, maximum, average and median results on the instance set for n , Q , r , n/K_{min} , as well as for the Gaps(%) and CPU time of ILS-SP, UHGS and BCP (root relaxation).

3.1. Heuristic solutions

We selected a pair of recent successful metaheuristics to illustrate the two main current types of approaches in the literature. We consider an efficient neighborhood-based method, the iterated local search based matheuristic algorithm (ILS-SP) of Subramanian, Uchoa, and Ochi (2013), and a recent population-based method, the unified hybrid genetic search (UHGS) of Vidal et al. (2012), Vidal, Crainic, Gendreau, and Prins (2014). These two methods rely extensively on local search to improve new solutions generated either by shaking or recombinations of parents. They also include specific strategies to explore new choices of customer-to-route assignments: ILS is coupled with an integer programming solver over a set partitioning (SP) formulation, which seeks to create new solutions based on known routes from past local optimums, while UHGS implements a continuous diversification procedure by modifying the objective during parents and survivors selection to promote not only good but also diverse solutions. Both methods are known to achieve very high quality results on the previous sets of CVRP instances as well as on various other vehicle routing variants.

The two methods have been run with the same parameter setting specified in the original papers. To produce solutions that can stand the test of time, and counting on the fact that more computing power will be surely available in the coming years, we selected a slightly larger termination criterion for UHGS by allowing up to 50,000 consecutive iterations without improvement. These tests have been conducted on a Xeon CPU with 3.07 gigahertz and 16 gigabytes of RAM, running under Oracle Linux Server 6.4. The average and best results on 50 runs are reported in the table, as well as the average computational time per instance.

From these tests, it appears that both ILS-SP and UHGS produce solutions of consistent quality, with an average gap of 0.52% and 0.19%, respectively, with respect to the best known solutions (BKS) ever found during all experiments. Considering the subset of instances that are solved to optimality (see Section 3.2), ILS-SP and UHGS achieve average gaps of 0.18% and 0.09%, respectively. A direct comparison of the best solutions found by each method provides the following score: UHGS solutions are better 58 times, ILS-SP solutions are better 9 times and there are 33 ties.

UHGS produces solutions of generally higher quality than ILS for a comparable amount of CPU time, except for some instances

Table 1

New set of benchmark instances: characteristics and results of current state-of-the-art algorithms (part I).

#	Name	Instance characteristics							ILS-SP			UHGS			BCP				BKS	
		n	Dep	Cust	Dem	Q	r	n/K_{min}	Avg	Best	$T(\text{min})$	Avg	Best	$T(\text{min})$	RLB	$RT(\text{min})$	Nds	$T(\text{min})$	Value	NV
1	X-n101-k25	100	R	RC (7)	1–100	206	4.0	4.0	27591.0	27,591	0.13	27591.0	27,591	1.43	27,591	0.1	1	0.1	27,591	26
2	X-n106-k14	105	E	C (3)	50–100	600	8.0	7.5	26375.9	26,362	2.01	26381.8	26,378	4.04	26,362	3.5	1	3.5	26,362	14
3	X-n110-k13	109	C	R	5–10	66	8.8	8.4	14971.0	14,971	0.20	14971.0	14,971	1.58	14,971	0.3	1	0.3	14,971	13
4	X-n115-k10	114	C	R	SL	169	12.5	11.4	12747.0	12,747	0.18	12747.0	12,747	1.81	12,747	2.1	1	2.1	12,747	10
5	X-n120-k6	119	E	RC (8)	U	21	21.8	19.8	13337.6	13,332	1.69	13332.0	13,332	2.31	13,234	2.2	63	88.1	13,332	6
6	X-n125-k30	124	R	C (5)	Q	188	4.2	4.1	55673.8	55,539	1.43	55542.1	55,539	2.66	55,539	2.5	1	2.5	55,539	30
7	X-n129-k18	128	E	RC (8)	1–10	39	7.4	7.1	28998.0	28,948	1.92	28948.5	28,940	2.71	28,897	1.3	3	2.5	28,940	18
8	X-n134-k13	133	R	C (4)	Q	643	10.4	10.2	10947.4	10,916	2.07	10934.9	10,916	3.32	10,840	8.3	2955	399.1	10,916	13
9	X-n139-k10	138	C	R	5–10	106	14.0	13.8	13603.1	13,590	1.60	13590.0	13,590	2.28	13,590	17.0	1	17.0	13,590	10
10	X-n143-k7	142	E	R	1–100	1190	22.6	20.3	15745.2	15,726	1.64	15700.2	15,700	3.10	15,634	20.9	1825	1553	15,700	7
11	X-n148-k46	147	R	RC (7)	1–10	18	3.2	3.2	43452.1	43,448	0.84	43448.0	43,448	3.18	43,448	0.3	1	0.3	43,448	47
12	X-n153-k22	152	C	C (3)	SL	144	7.1	6.9	21400.0	21,340	0.49	21226.3	21,220	5.47	21,140	4.7	143	37.7	21,220	23
13	X-n157-k13	156	R	C (3)	U	12	12.0	12.0	16876.0	16,876	0.76	16876.0	16,876	3.19	16,876	1.0	1	1.0	16,876	13
14	X-n162-k11	161	C	RC (8)	50–100	1174	15.5	14.6	14160.1	14,138	0.54	14141.3	14,138	3.32	14,053	35.6	101	187.0	14,138	11
15	X-n167-k10	166	E	R	5–10	133	17.8	16.6	20608.7	20,562	0.86	20563.2	20,557	3.73	20,476	8.5	85	1024	20,557	10
16	X-n172-k51	171	C	RC (5)	Q	161	3.4	3.4	45616.1	45,607	0.64	45607.0	45,607	3.83	45,549	1.5	3	3.8	45,607	53
17	X-n176-k26	175	E	R	SL	142	6.8	6.7	48249.8	48,140	1.11	47957.2	47,812	7.56	47,721	2.4	43	9.2	47,812	26
18	X-n181-k23	180	R	C (6)	U	8	8.4	7.8	25571.5	25,569	1.59	25591.1	25,569	6.28	25,511	1.9	73	18.2	25,569	23
19	X-n186-k15	185	R	R	50–100	974	13.0	12.3	24186.0	24,145	1.72	24147.2	24,145	5.92	23,980	26.8	1211	7305	24,145	15
20	X-n190-k8	189	E	C (3)	1–10	138	25.0	23.6	17143.1	17,085	2.10	16987.9	16,980	12.08	16,939	22.8	993	16146	16,980	8
21	X-n195-k51	194	C	RC (5)	1–100	181	3.8	3.8	44234.3	44,225	0.87	44244.1	44,225	6.10	44,225	2.4	1	2.4	44,225	53
22	X-n200-k36	199	R	C (8)	Q	402	5.6	5.5	58697.2	58,626	7.48	58626.4	58,578	7.97	58,455	1.7	1469	901.7	58,578	36
23	X-n204-k19	203	C	RC (6)	50–100	836	11.2	10.7	19625.2	19,570	1.08	19571.5	19,565	5.35	19,484	75.4	85	501.6	19,565	19
24	X-n209-k16	208	E	R	5–10	101	13.5	13.0	30765.4	30,667	3.80	30680.4	30,656	8.62	30,480	3.2	603	1303	30,656	16
25	X-n214-k11	213	C	C (4)	1–100	944	19.4	19.4	11126.9	10,985	2.26	10877.4	10,856	10.22	10,809	264.9			10,856	11
26	X-n219-k73	218	E	R	U	3	3.6	3.0	117595.0	117,595	0.85	117604.9	117,595	7.73	117,595	0.5	1	0.5	117,595	73
27	X-n223-k34	222	R	RC (5)	1–10	37	6.5	6.5	40533.5	40,471	8.48	40499.0	40,437	8.26	40,311	13.2	343	303.5	40,437	34
28	X-n228-k23	227	R	C (8)	SL	154	10.0	9.9	25795.8	25,743	2.40	25779.3	25,742	9.80	25,657	15.0	163	252.3	25,742	23
29	X-n233-k16	232	C	RC (7)	Q	631	14.5	14.5	19336.7	19,266	3.01	19288.4	19,230	6.84	19,070	133.0			19,230	17
30	X-n237-k14	236	E	R	U	18	18.6	16.9	27078.8	27,042	3.46	27067.3	27,042	8.90	26,930	2.0	2695	1398	27,042	14
31	X-n242-k48	241	E	R	1–10	28	5.0	5.0	82874.2	82,774	17.83	82948.7	82,804	12.42	82,589	2.0	2661	819.0	82,751	48
32	X-n247-k50	246	C	C (4)	SL	134	5.3	4.9	37507.2	37,289	2.06	37284.4	37,274	20.41	37,256	3.7	7	9.5	37,274	51
33	X-n251-k28	250	R	RC (3)	5–10	69	9.2	8.9	38840.0	38,727	10.77	38796.4	38,699	11.69	38,473	2.1	2531	4767	38,684	28
34	X-n256-k16	255	C	C (8)	50–100	1225	16.0	15.9	18883.9	18,880	2.02	18880.0	18,880	6.52	18,826	201.5	25	1255	18,880	17
35	X-n261-k13	260	E	R	1–100	1081	21.0	20.0	26869.0	26,706	6.67	26629.6	26,558	12.67	26,407	373.7			26,558	13

Table 2

New set of benchmark instances: characteristics and results of current state-of-the-art algorithms (part II).

#	Name	Instance characteristics							ILS-SP			UHGS			BCP				BKS	
		<i>n</i>	Dep	Cust	Dem	<i>Q</i>	<i>r</i>	<i>n/K_{min}</i>	Avg	Best	<i>T</i> (min)	Avg	Best	<i>T</i> (min)	RLB	<i>RT</i> (min)	Nds	<i>T</i> (min)	Value	NV
36	X-n266-k58	265	R	RC (6)	5–10	35	4.6	4.6	75563.3	75,478	10.03	75759.3	75,517	21.36	75,350	9.5	185	150.1	75,478	58
37	X-n270-k35	269	C	RC (5)	50–100	585	7.7	7.7	35363.4	35,324	9.07	35367.2	35,303	11.25	35,156	14.5	389	3422	35,291	36
38	X-n275-k28	274	R	C (3)	U	10	10.8	9.8	21256.0	21,245	3.59	21280.6	21,245	12.04	21,245	3.7	1	3.7	21,245	28
39	X-n280-k17	279	E	R	SL	192	16.5	16.4	33769.4	33,624	9.62	33605.8	33,505	19.09	33,286	87.7			33,503	17
40	X-n284-k15	283	R	C (8)	1–10	109	20.2	18.9	20448.5	20,295	8.64	20286.4	20,227	19.91	20,139	49.4			20,226	15
41	X-n289-k60	288	E	RC (7)	Q	267	4.8	4.8	95450.6	95,315	16.11	95469.5	95,244	21.28	94,928	25.1	8220	24708	95,151	61
42	X-n294-k50	293	C	R	1–100	285	5.9	5.9	47254.7	47,190	12.42	47259.0	47,171	14.70	46,911	26.7			47,167	51
43	X-n298-k31	297	R	R	1–10	55	9.6	9.6	34356.0	34,239	6.92	34292.1	34,231	10.93	34,105	1.8	195	531.1	34,231	31
44	X-n303-k21	302	C	C (8)	1–100	794	15.0	14.4	21895.8	21,812	14.15	21850.9	21,748	17.28	21,546	481.0			21,744	21
45	X-n308-k13	307	E	RC (6)	SL	246	24.2	23.6	26101.1	25,901	9.53	25895.4	25,859	15.31	25,587	247.7			25,859	13
46	X-n313-k71	312	R	RC (3)	Q	248	4.4	4.4	94297.3	94,192	17.50	94265.2	94,093	22.41	93,851	11.6			94,044	72
47	X-n317-k53	316	E	C (4)	U	6	6.2	6.0	78356.0	78,355	8.56	78387.8	78,355	22.37	78,334	1.7	5	4.6	78,355	53
48	X-n322-k28	321	C	R	50–100	868	11.6	11.5	29991.3	29,877	14.68	29956.1	29,870	15.16	29,722	547.2			29,866	28
49	X-n327-k20	326	R	RC (7)	5–10	128	17.0	16.3	27812.4	27,599	19.13	27628.2	27,564	18.19	27,378	221.9			27,556	20
50	X-n331-k15	330	E	R	U	23	23.4	22.0	31235.5	31,105	15.70	31159.6	31,103	24.43	31,027	25.7			31,103	15
51	X-n336-k84	335	E	R	Q	203	4.0	4.0	139461.0	139,197	21.41	139534.9	139,210	37.96	138,706	6.6			139,197	86
52	X-n344-k43	343	C	RC (7)	5–10	61	8.0	8.0	42284.0	42,146	22.58	42208.8	42,099	21.67	41,881	20.4			42,099	43
53	X-n351-k40	350	C	C (3)	1–100	436	8.8	8.8	26150.3	26,021	25.21	26014.0	25,946	33.73	25,809	170.4			25,946	41
54	X-n359-k29	358	E	RC (7)	1–10	68	12.5	12.3	52076.5	51,706	48.86	51721.7	51,509	34.85	51,381	82.3			51,509	29
55	X-n367-k17	366	R	C (4)	SL	218	21.8	21.5	23003.2	22,902	13.13	22838.4	22,814	22.02	22,747	247.9			22,814	17
56	X-n376-k94	375	E	R	U	4	4.2	4.0	147713.0	147,713	7.10	147750.2	147,717	28.26	147,713	3.3	1	3.3	147,713	94
57	X-n384-k52	383	R	R	50–100	564	7.4	7.4	66372.5	66,116	34.47	66270.2	66,081	40.20	65,681	256.7			66,081	53
58	X-n393-k38	392	C	RC (5)	5–10	78	10.4	10.3	38457.4	38,298	20.82	38374.9	38,269	28.65	38,167	46.7			38,269	38
59	X-n401-k29	400	E	C (6)	Q	745	14.0	13.8	66715.1	66,453	60.36	66365.4	66,243	49.52	65,971	318.1			66,243	29
60	X-n411-k19	410	R	C (5)	SL	216	22.6	21.6	19954.9	19,792	23.76	19743.8	19,718	34.71	19,640	433.9			19,718	19
61	X-n420-k130	419	C	RC (3)	1–10	18	3.2	3.2	107838.0	107,798	22.19	107924.1	107,798	53.19	107,704	5.4	169	115.0	107,798	130
62	X-n429-k61	428	R	R	50–100	536	7.1	7.0	65746.6	65,563	38.22	65648.5	65,501	41.45	64,930	31.6			65,501	62
63	X-n439-k37	438	C	RC (8)	U	12	12.0	11.8	36441.6	36,395	39.63	36451.1	36,395	34.55	36,289	20.0	2845	33615	36,391	37
64	X-n449-k29	448	E	R	1–100	777	15.5	15.4	56204.9	55,761	59.94	55553.1	55,378	64.92	54,928	804.9			55,358	29
65	X-n459-k26	458	C	C (4)	Q	1106	17.8	17.6	24462.4	24,209	60.59	24272.6	24,181	42.80	23,931	327.1			24,181	26
66	X-n469-k138	468	E	R	50–100	256	3.4	3.4	222182.0	221,909	36.32	222617.1	222,070	86.65	221,429	10.9			221,909	140
67	X-n480-k70	479	R	C (8)	5–10	52	6.8	6.8	89871.2	89,694	50.40	89760.1	89,535	66.96	89,235	59.5			89,535	70
68	X-n491-k59	490	R	RC (6)	1–100	428	8.4	8.3	67226.7	66,965	52.23	66898.0	66,633	71.94	66,263	418.1			66,633	60
69	X-n502-k39	501	E	C (3)	U	13	13.0	12.8	69346.8	69,284	80.75	69328.8	69,253	63.61	69,120	16.5			69,253	39
70	X-n513-k21	512	C	RC (4)	1–10	142	25.0	24.4	24434.0	24,332	35.04	24296.6	24,201	33.09	24,053	262.1			24,201	21

Table 3

New set of benchmark instances: characteristics and results of current state-of-the-art algorithms (part III).

#	Name	Instance Characteristics							ILS-SP			UHGS			BCP			BKS		
		<i>n</i>	Dep	Cust	Dem	<i>Q</i>	<i>r</i>	<i>n/K_{min}</i>	Avg	Best	<i>T</i> (min)	Avg	Best	<i>T</i> (min)	RLB	<i>RT</i> (min)	Nds	<i>T</i> (min)	Value	NV
71	X-n524-k153	523	R	R	SL	125	3.8	3.4	155005.0	154,709	27.27	154979.5	154,774	80.70	154,533	12.5	381	212.1	154,594	155
72	X-n536-k96	535	C	C (7)	Q	371	5.6	5.6	95700.7	95,524	62.07	95330.6	95,122	107.53	94,409	47.9			95,122	97
73	X-n548-k50	547	E	R	U	11	11.2	10.9	86874.1	86,710	63.95	86998.5	86,822	84.24	86,604	16.1			86,710	50
74	X-n561-k42	560	C	RC (7)	1–10	74	13.5	13.3	43131.3	42,952	68.86	42866.4	42,756	60.60	42,495	302.3			42,756	42
75	X-n573-k30	572	E	C (3)	SL	210	19.4	19.1	51173.0	51,092	112.03	50915.1	50,780	188.15	50,575	1306.9			50,780	30
76	X-n586-k159	585	R	RC (4)	5–10	28	3.6	3.7	190919.0	190,612	78.54	190838.0	190,543	175.29	189,950	7.1			190,543	159
77	X-n599-k92	598	R	R	50–100	487	6.5	6.5	109384.0	109,056	72.96	109064.2	108,813	125.91	108,000	264.3			108,813	94
78	X-n613-k62	612	C	R	1–100	523	10.0	9.9	60444.2	60,229	74.80	59960.0	59,778	117.31	59,323	1429.8			59,778	62
79	X-n627-k43	626	E	C (5)	5–10	110	14.5	14.6	62905.6	62,783	162.67	62524.1	62,366	239.68	62,018	600.2			62,366	43
80	X-n641-k35	640	E	RC (8)	50–100	1381	18.6	18.3	64606.1	64,462	140.42	64192.0	63,839	158.81	63,228	401.0			63,839	35
81	X-n655-k131	654	C	C (4)	U	5	5.0	5.0	106782.0	106,780	47.24	106899.1	106,829	150.48	106,766	8.5	21	41.5	106,780	131
82	X-n670-k130	669	R	R	SL	129	5.3	5.1	147676.0	147,045	61.24	147222.7	146,705	264.10	146,211	103.1			146,705	134
83	X-n685-k75	684	C	RC (6)	Q	408	9.2	9.1	68988.2	68,646	73.85	68654.1	68,425	156.71	67,925	1643.2			68,425	75
84	X-n701-k44	700	E	RC (7)	1–10	87	16.0	15.9	83042.2	82,888	210.08	82487.4	82,293	253.17	81,694	680.6			82,292	44
85	X-n716-k35	715	R	C (3)	1–100	1007	21.0	20.4	44171.6	44,021	225.79	43641.4	43,525	264.28	43,113	419.5			43,525	35
86	X-n733-k159	732	C	R	1–10	25	4.6	4.6	137045.0	136,832	111.56	136587.6	136,366	244.53	135,748	57.9			136,366	160
87	X-n749-k98	748	R	C (8)	1–100	396	7.7	7.6	78275.9	77,952	127.24	77864.9	77,715	313.88	76,924	259.2			77,700	98
88	X-n766-k71	765	E	RC (7)	SL	166	10.8	10.8	115738.0	115,443	242.11	115147.9	114,683	382.99	114,108	262.8			114,683	71
89	X-n783-k48	782	R	R	Q	832	16.5	16.3	73722.9	73,447	235.48	73009.6	72,781	269.70	71,728	332.8			72,727	48
90	X-n801-k40	800	E	R	U	20	20.2	20.0	74005.7	73,830	432.64	73731.0	73,587	289.24	73,124	258.0			73,587	40
91	X-n819-k171	818	C	C (6)	50–100	358	4.8	4.8	159425.0	159,164	148.91	158899.3	158,611	374.28	157,627	257.2			158,611	173
92	X-n837-k142	836	R	RC (7)	5–10	44	5.9	5.9	195027.0	194,804	173.17	194476.5	194,266	463.36	193,245	253.7			194,266	142
93	X-n856-k95	855	C	RC (3)	U	9	9.6	9.0	89277.6	89,060	153.65	89238.7	89,118	288.43	88,839	43.6			89,060	95
94	X-n876-k59	875	E	C (5)	1–100	764	15.0	14.8	100417.0	100,177	409.31	99884.1	99,715	495.38	98,880	433.2			99,715	59
95	X-n895-k37	894	R	R	50–100	1816	24.2	24.2	54958.5	54,713	410.17	54439.8	54,172	321.89	53,147	984.5			54,172	38
96	X-n916-k207	915	E	RC (6)	5–10	33	4.4	4.4	330948.0	330,639	226.08	330198.3	329,836	560.81	328,588	97.1			329,836	208
97	X-n936-k151	935	C	R	SL	138	6.2	6.2	134530.0	133,592	202.50	133512.9	133,140	531.50	132,496	395.8			133,105	159
98	X-n957-k87	956	R	RC (4)	U	11	11.6	11.0	85936.6	85,697	311.20	85822.6	85,672	432.90	85,328	257.1			85,672	87
99	X-n979-k58	978	E	C (6)	Q	998	17.0	16.9	120253.0	119,994	687.22	119502.1	11,9194	553.96	118,399	937.5			119,194	58
100	X-n1001-k43	1000	R	R –	1–10	131	23.4	23.3	73985.4	73,776	792.75	72956.0	72,742	549.03	71,812	308.8			72,742	43
Min		100				3	3.2	3.0	0.00%	0.00%	0.13	0.00%	0.00%	1.43	0.00%	0.1				
Max		1000				1816	25.0	24.4	2.50%	1.42%	792.75	0.55%	0.13%	560.81	1.89%	1643.2				
Avg.		412.2				324.6	11.4	11.1	0.52%	0.25%	71.71	0.19%	0.01%	98.79	0.44%	189.4				
Median		333				149	10.2	9.9	0.38%	0.10%	17.67	0.20%	0.00%	22.39	0.40%	33.6				

containing few customers per route. For this type of problems, the set partitioning solver can produce new high-quality solutions from existing routes in a very efficient manner, while generating new structurally different solutions with other choices of assignments is more challenging for local searches and even for randomized crossovers. On the other hand, for problems with a large number of customers per route, the hybrid ILS exhibits a slower convergence and generally leads to solutions of lower quality.

Overall, these experiments show that some state-of-the-art methods may perform well on different classes of instances. Therefore, a promising research path would involve the extension of these methods and/or further hybridizations to cover all problems in the best possible way.

3.2. Exact solutions

The Branch-Cut-and-Price (BCP) in [Pecin et al. \(2014\)](#) is a complex algorithm that combines and improves several ideas proposed by other authors ([Baldacci et al., 2008](#); [Baldacci et al., 2011](#); [Contardo & Martinelli, 2014](#); [Røpke, 2012](#)), as also detailed in [Poggi and Uchoa \(2014\)](#). The algorithm was run in a single core of an Intel i7-3960X 3.30 gigahertz processor with 64 gigabytes RAM. The value of the best solution found by the previously mentioned heuristics was inputted as an initial upper bound, and for a fair computational comparison, we increased this bound by 1 unit, in such a way that the BCP always has to find by itself a feasible solution with value at least as good as the upper bound to prove optimality.

The BCP could solve 43 out of the 100 new instances to optimality. For those instances we report the root node lower bound, the time to obtain that bound, the number of nodes in the search tree and the total time. For the remaining 57 instances we only report a lower bound and the corresponding time.

Most instances with up to 275 customers was solved by the BCP. For the 38 instances in that range, it only failed when the routes are very long (X-n214-k11, X-n233-k16, and X-n261-k13). On the other hand, only 8 larger instances could be solved (X-n289-k60, X-n298-k31, X-n317-k53, X-376-k94, X-n420-k130, X-n439-k37, X-n524-k153, and X-nX655-k131). Those more favorable instances have short routes or small values of Q .

The BCP attested the general good quality of the solutions that can be found by current heuristics. In 96 out of the 100 instances, the best solution found (provided either by ILS-SP or UHGS) is less than 1% above the computed lower bounds.

4. Extended benchmark

The previous set of 100 benchmark instances was designed for evaluating the performance of different algorithms, allowing for detailed comparisons of heuristic or exact methods and solution reports within moderate CPU time. However, since there are five main instance attributes that can assume several possible values (as described in [Section 2](#)), the number of instances presented in [Section 3](#) is statistically insufficient to assess the impact of each of the attributes on the performance of the algorithms. To go one step further, we created a second set of instances using a one-factor-at-a-time (OFAT) approach ([Frey & Wang, 2006](#)). Such an approach is usually less statistically powerful than other more advanced methods for design of experiments, but it has the merit of being intuitive and easy to interpret. We thus defined a *standard* instance configuration with $n = 200$, random depot positioning, random-clustered customer positioning, small values and small CV demand distribution (5–10), and average route size in [9, 11]. Those factors were varied, one at a time, as described in [Table 4](#), to obtain 29

Table 4

Attribute values of the extended benchmark, one-factor-at-a-time approach. Values of the standard configuration are underlined.

n	Depot	Customer	Demand	r
100	C	R	U	[3,5]
125	E	C	1–10	[6,8]
150	<u>R</u>	<u>RC</u>	<u>5–10</u>	<u>[9,11]</u>
175			1–100	[12,14]
<u>200</u>			50–100	[15,16]
225			Q	[18,19]
250			SL	[21,22]
300				
350				
400				
450				
500				
550				
600				

other configurations. The first group of 14 non-standard configurations only differs from the standard one by changing the value of n , the second group of two configurations differs from the standard only in the depot positioning, and so on. We generated 20 random instances per configuration. Note that these instances, a total of 600, are provided not with the aim to *compete* for producing better solutions, but as a tool for further analysis of algorithms.

4.1. Analysis of heuristics

Each of the two metaheuristics (UHGS and ILS) was run ten times on each instance using the same parameter setting as in [Section 3](#). We collected the average solution of these runs to measure the impact of the instance characteristics on the performance of each method in terms of the percentage gap from the best known solutions (BKS), extracted from all heuristic and exact runs. [Fig. 1](#) reports these results in the form of a boxplot for each method and instance configuration. As such, each boxplot correspond to 20 different instances of the same type, and each measure is an average of ten runs.

Firstly, the performance of UHGS and ILS behaves differently as a function of the average number of customers per route. The hybrid ILS takes a large advantage of short routes, due to the use of the set partitioning formulation, but its performance significantly worsens for problems with longer routes. In contrast, the performance of UHGS generally improves as the routes grow larger.

Problem size is, as expected, a key factor for heuristic performance. Small problem instances with less than 150 customers are almost systematically solved to optimality by UHGS, while large-size problem instances can lead to gaps of up to 0.64% and 1.05%, for UHGS and ILS, respectively. The average percentage gap of UHGS tends to first increase with the number of customers, from 0.01% on problems with 100 deliveries to 0.28% on problems with 350 deliveries, and then reach a *performance plateau*, characterized by a smaller variation of the gap measure, located between 0.27% and 0.34% for larger instances with 400–600 customers.

With respect to the depot positioning, both heuristics reveal a similar behavior in our tests. Instances with a centered depot are more easy than those with a depot located in a corner, as this latter feature leads to more tours oriented towards the same area, and thus more opportunities and complexity for intra-route improvements. Moreover, instances with clustered customers are far more easily solved by heuristics, as this feature tends to limit the pool of desirable routes and reduce route overlaps in different regions.

Finally, various customer-demands configurations have been tested. Instances with unitary demands tend to be easier for the




Fig. 1. Impact of the instance parameters (problem size, route size, depot positioning, customer demands and distribution) on heuristic resolution performance.

Table 5
BCP median times for each configuration.

	Hours	Days	Ratio		Hours	Days	Ratio
<i>n</i>				Demand			
100	0.26	0.01	0.03	U	0.19	0.01	0.02
125	0.2	0.01	0.02	1–10	4.41	0.18	0.45
150	1.36	0.06	0.14	5–10	9.73	0.41	1
175	2.99	0.12	0.31	1–100	99.51	4.15	10.23
200	9.73	0.41	1	50–100	183.01	7.63	18.81
225	40.85	1.7	4.2	Q	>312	>13	>32.07
250	34.77	1.45	3.57	SL	286.5	11.94	29.45
Dep				<i>r</i>			
C	1.42	0.06	0.15	[3,5]	1.37	0.06	0.14
E	53.15	2.21	5.46	[6,8]	2.03	0.08	0.21
R	9.73	0.41	1	[9,11]	9.73	0.41	1
				[12,14]	24.14	1.01	2.48
Cust				[15,16]	86.67	3.61	8.91
R	9.92	0.41	1.02	[17,18]	139.9	5.83	14.38
C	7.1	0.3	0.73	[21,23]	143.32	5.97	14.73
RC	9.73	0.41	1				

hybrid ILS. A higher coefficient of variation for the customer demands renders the problem instances moderately more difficult. The assignment of customers to routes is indeed more complex with customer demands of different scale. Finally, ILS and UHGS exhibit different behavior in the presence of customer demands which are dependent on the region (quadrant configuration), or with a few large deliveries. For these problems, the performance of ILS tends to deteriorate, again due to the fact that the set partitioning models are harder to solve.

4.2. Analysis of the exact method

The Branch-Cut-and-Price (BCP) was executed in the same machine used in the previous tests, but with 12 instances being run in parallel, one in each available core. As larger values of n makes the CPU times of the exact algorithm prohibitively high, we limited our experiments to the 460 instances of the extended benchmark with less than 250 customers. The first observation was that the CPU times of the BCP can vary a lot even on instances with the same configuration. For example, consider the 20 instances of the standard configuration. While 8 instances were solved in less than 5 hours, the two hardest instances took 7 and 11 days of CPU time. The average time of 1.52 days was heavily driven by those two instances. In those conditions, we believe that the median time (0.41 days for the standard configuration) is a more representative measure of the relative difficulty of each configuration. The use of medians is also due to a practical reason. The calculation of the averages for a given configuration would require solving all the 20 instances in it, and this could be extremely time consuming in some cases. On the other hand, only 11 instances of a configuration need to be solved to optimality for calculating the medians. In fact, we could abort the long runs of several instances after the median of their configurations was already determined. The median CPU times in Table 5 are presented both in hours and days. We also report the ratio between the CPU time of a given configuration and the CPU time of the standard configuration.

As expected, the CPU times grow nearly exponentially with n in the interval 100–250. A linear regression of the CPU time T in hours, as a function of the number of customers leads to $\log(T) = 0.0169n - 2.4605$. Extrapolating this trend for $n = 300$ gives a predicted median time of 17 days, justifying the decision of not running instances with n larger than 250.

The influence of the depot positioning on the CPU time is very evident. The median time for instances with the depot in the corner was 37 times larger than those with the depot in the center. Instances with random depot location have intermediate difficulty.

On the other hand, customer positioning does not seem to considerably affect the performance of BCP, and clustered instances appear to be just a bit easier.

Demand distribution has a very significant influence on the CPU time. Unitary demands make instances much easier. It can be noted that unitary demand instances with $n = 200$ are better solved than instances with $n = 100$ and demands in [5–10]. Instances with larger demand values ([1–100] and [50–100]) are harder than instances with small values ([1–10] and [5–10]). This is expected because the pricing phase of BCP uses a dynamic programming with a pseudo-polynomial complexity. Demands with large CV ([1–10] and [1–100]) turned out to be easier than their counterparts with small CV ([5–10] and [50–100], respectively). Instances with few very large demands and many small demands were difficult to solve, with a median time of almost 12 days. The hardest of all categories was formed by the instances with demand dependent on the quadrant. Only 8 such instances could be solved in less than 9 days, the runs of the remaining 12 instances were aborted with at least 13 days of CPU time, so the median time is larger than 13 days.

The experiments finally confirmed that the average size of the routes has a large impact on the CPU times of the BCP.

5. The CVRPLIB web site

The typical instance repository of today is a web page that allows for downloading the instance files and includes additional textual information, such as the file format description, instance source and best known/optimal solution values, among others. The CVRPLIB web page (<http://vrp.atd-lab.inf.puc-rio.br/>), containing the new instances and all the previous CVRP instances described in the appendix, is more sophisticated:

- Its core is a full-fledged database containing, for each instance: (i) its actual data, n , Q , customer coordinates and demands, (ii) its best known/optimal solution, represented as set of routes, (iii) a miscellanea of additional information, e.g., the original source, the values of the attributes used in its generation (for the new instances), and some comments about remarkable characteristics. The visible pieces of information that appear in the web site are produced by queries. The objective of that design is having a larger degree of consistency and flexibility in the maintenance of the page.
- The best known solutions are automatically verified, their feasibility is checked and their costs are calculated from the original instance data. This eliminates the possibility that false best known solutions appear due to typos or misunderstandings about the conventions.
- Instances with Euclidean distances (the vast majority) can be depicted graphically, along with their best known/optimal solutions.

6. Conclusions

We hope that the proposed set of instances will contribute to spark new interest on the “classic” CVRP, helping to test new algorithmic developments in the coming years. In particular, by proposing a common set of benchmark instances for both heuristic and exact methods, this work contributes to end a bizarre situation where minor convention details were sufficient for isolating both communities. For example, most of the CVRP instances in Table A.12 could have their best known solutions proven to be optimal by the algorithms in Fukasawa et al. (2006), Baldacci et al. (2008), Baldacci et al. (2011), Røpke (2012), Contardo and Martinelli (2014). But none of those authors effectively changed a few lines of code to adapt for the different conventions. This

instance-induced isolation was harmful and prevented a more effective cross-fertilization between exact and heuristic methods. In fact, this paper has shown that it is already possible to perform direct comparisons of heuristic results with good lower bounds or even proven optimal solutions on fairly large-sized instances.

A major concern in the design of the new benchmark set was problem diversity. By having instances covering a wider set of characteristics, the benchmark will favor the development of more flexible methods (perhaps hybrids), that will eventually be applied to real-life situations with lesser risks of failure due to an uncommon instance type.

Finally, we point out that the proposed CVRP instances can be extended to other classical VRP variants as needed, by providing additional data fields like duration constraints, time windows, or heterogeneous fleet specifications.

Acknowledgments

The authors would like to thank Teobaldo Bulhões Jr. for computing the values of K_{min} for the new set of proposed instances and Ivan Xavier de Lima for his help on creating the CVRPLIB web page.

Appendix A. Current benchmark instances

In this appendix we provide detailed information about the existing benchmark instances. In particular, we tried to track back the generation process of all these instances.

Table A.6 presents data regarding the E series. Although they are usually attributed to Christofides and Eilon (1969), some instances actually come from Dantzig and Ramser (1959) and from Gaskell (1967), and some are modifications later suggested by Gillett and Miller (1974). As usual in the literature on exact methods, the naming of the instances reflects the number of points (including the depot) and the fixed number of routes. For example, E-n101-k8 is an instance with 100 customers and a requirement of 8 routes. Columns in Table A.6 include the value of Q , the tightness (the ratio between the sum of all demands and KQ , the total capacity available in the fixed number of routes) and the optimal solution value.

Table A.7 presents information about the M series and explains how each instance was generated. Instances M-n200-k16 and M-n200-k17 only differ by the required number of routes. In fact, M-n200-k17 is the only ABEFMP instance where the fixed number of routes does not match the minimum possible. This additional in-

Table A.6
Instances of the set E.

Instance	Q	Tightness	Opt	Original source
E-n13-k4 ¹	6000	0.76	247	Dantzig and Ramser (1959)
E-n22-k4 ¹	6000	0.94	375	Gaskell (1967)
E-n23-k3 ¹	4500	0.75	569	Gaskell (1967)
E-n30-k3 ¹	4500	0.94	534	Gaskell (1967)
E-n31-k7 ²	140	0.92	379	Clarke and Wright (1964)
E-n33-k4 ¹	8000	0.92	835	Gaskell (1967)
E-n51-k5 ³	160	0.97	521	Christofides and Eilon (1969)
E-n76-k7 ⁴	220	0.89	682	Gillett and Miller (1974)
E-n76-k8 ⁴	180	0.95	735	Gillett and Miller (1974)
E-n76-k10 ³	140	0.97	830	Christofides and Eilon (1969)
E-n76-k14 ⁴	100	0.97	1021	Gillett and Miller (1974)
E-n101-k8 ³	200	0.91	815	Christofides and Eilon (1969)
E-n101-k14 ⁵	112	0.93	1067	Gillett and Miller (1974)

¹ No description about the generation.

² Example involving UK cities with customers located far from the depot.

³ Locations generated at random from a uniform distribution.

⁴ Instance E-n76-k10 with modified capacity.

⁵ Instance E-n101-k8 with modified capacity.

Table A.7

Instances of the set M, original source: Christofides et al. (1979).

Instance	Q	Tightness	Opt
M-n101-k10 ¹	200	0.91	820
M-n121-k7 ¹	200	0.98	1034
M-n151-k12 ²	200	0.93	1015
M-n200-k16 ³	200	1.00	1274
M-n200-k17 ³	200	0.94	1275

¹ Customers were grouped into clusters as an attempt to represent practical cases.

² Generated by adding customers from E-n51-k5 and E-n101-k8 and using the depot and capacity from E-n101-k8.

³ Generated by adding customers from M-n151-k12 and the first 49 customers from E-n76-k10 and using the depot and capacity from M-n151-k12.

Table A.8

Instances of the set F, original source: Fisher (1994).

Instance	Q	Tightness	Opt
F-n45-k4 ¹	2010	0.90	724
F-n72-k4 ²	30000	0.96	237
F-n135-k7 ¹	2210	0.95	1162

¹ From a day of grocery deliveries from the Peterboro (Ontario terminal) of National Grocers Limited.

² Data obtained from Exxon associated to the delivery of tires, batteries and accessories to gasoline service stations.

Table A.9

Instances of the set A, original source: Augerat et al. (1995)¹.

Instance	Q	Tightness	Opt
A-n32-k5	100	0.82	784
A-n33-k5	100	0.89	661
A-n33-k6	100	0.90	742
A-n34-k5	100	0.92	778
A-n36-k5	100	0.88	799
A-n37-k5	100	0.81	669
A-n37-k6	100	0.95	949
A-n38-k5	100	0.96	730
A-n39-k5	100	0.95	822
A-n39-k6	100	0.88	831
A-n44-k6	100	0.95	937
A-n45-k6	100	0.99	944
A-n45-k7	100	0.91	1146
A-n46-k7	100	0.86	914
A-n48-k7	100	0.89	1073
A-n53-k7	100	0.95	1010
A-n54-k7	100	0.96	1167
A-n55-k9	100	0.93	1073
A-n60-k9	100	0.92	1354
A-n61-k9	100	0.98	1034
A-n62-k8	100	0.92	1288
A-n63-k9	100	0.97	1616
A-n63-k10	100	0.93	1314
A-n64-k9	100	0.94	1401
A-n65-k9	100	0.97	1174
A-n69-k9	100	0.94	1159
A-n80-k10	100	0.94	1763

¹ Coordinates are random points in a $[0, 100] \times [0, 100]$ grid. Demands are picked from a uniform distribution $U(1,30)$, however $n/10$ of those demands are multiplied by 3.

stance was created because M-n200-k16 has a tightness so close to 1 (0.995625) that finding good feasible solutions for it was very difficult. Surprisingly, it was recently discovered that the optimal solution of M-n200-k16 costs less than the optimal solution of M-n200-k17 (Pecin et al., 2014).

Table A.8 presents the three real-world instances that compose the F series. Tables A.9 and A.10 correspond to the A and B series, respectively. While in the A series the customers and the depot are randomly positioned; they are clustered in the B series. The instances from series P were generated by taking some instances

Table A.10Instances of the set B, original source: Augerat et al. (1995)¹.

Instance	Q	Tightness	Opt
B-n31-k5	100	0.82	672
B-n34-k5	100	0.91	788
B-n35-k5	100	0.87	955
B-n38-k6	100	0.85	805
B-n39-k5	100	0.88	549
B-n41-k6	100	0.95	829
B-n43-k6	100	0.87	742
B-n44-k7	100	0.92	909
B-n45-k5	100	0.97	751
B-n45-k6	100	0.99	678
B-n50-k7	100	0.87	741
B-n50-k8	100	0.92	1312
B-n51-k7	100	0.98	1032
B-n52-k7	100	0.87	747
B-n56-k7	100	0.88	707
B-n57-k7	100	1.00	1153
B-n57-k9	100	0.89	1598
B-n63-k10	100	0.92	1496
B-n64-k9	100	0.98	861
B-n66-k9	100	0.96	1316
B-n67-k10	100	0.91	1032
B-n68-k9	100	0.93	1272
B-n78-k10	100	0.94	1221

¹ Coordinates are points in a $[0, 100] \times [0, 100]$ grid, chosen in order to create NC clusters. In all instances, $K \leq NC - 1$. Demands are picked from a uniform distribution $U(1,30)$, however $n/10$ of those demands are multiplied by 3.

Table A.11Instances of the set P, original source: Augerat et al. (1995)¹.

Instance	Q	Tightness	Opt
P-n16-k8	35	0.88	450
P-n19-k2	160	0.97	212
P-n20-k2	160	0.97	216
P-n21-k2	160	0.93	211
P-n22-k2	160	0.96	216
P-n22-k8	3000	0.94	603
P-n23-k8	40	0.98	529
P-n40-k5	140	0.88	458
P-n45-k5	150	0.92	510
P-n50-k7	150	0.91	554
P-n50-k8	120	0.99	631
P-n50-k10	100	0.95	696
P-n51-k10	80	0.97	741
P-n55-k7	170	0.88	568
P-n55-k8	160	0.81	588
P-n55-k10	115	0.91	694
P-n55-k15	70	0.99	989
P-n60-k10	120	0.95	744
P-n60-k15	80	0.95	968
P-n65-k10	130	0.94	792
P-n70-k10	135	0.97	827
P-n76-k4	350	0.97	593
P-n76-k5	280	0.97	627
P-n101-k4	400	0.91	681

¹ Modifications in the capacity of some instances from A, B and E series. Required number of routes are adjusted accordingly.

from the A, B and E series and changing their capacities. Consequently, the required number of routes also changes. For example, P-n101-k4 was obtained from E-n101-k8 by doubling the capacity.

Christofides et al. (1979) defined the benchmark set shown in Table A.12. Instances CMT1, CMT2, CMT3, CMT4, CMT5, CMT11, and CMT12 correspond to instances E-n51-k5, E-n76-k10, E-n101-k8, M-n151-k12, M-n200-k16, M-n121-k7, and M-n100-k10, respectively. The only difference are the conventions: the Euclidean distances are represented with full computer precision (without rounding), and the number of routes is not fixed. This set also contains instances for the duration constrained CVRP, obtained from the previous CVRP instances by adding maximum route duration (MD) and service time (ST) values. These values are also reported

Table A.12

Instances of Christofides et al. (1979).

Instance	n	Q	MD	ST	BKS	K_{BKS}
CMT1 ¹	50	160	∞	0	524.61*	5
CMT2 ¹	75	140	∞	0	835.26*	10
CMT3 ¹	100	200	∞	0	826.14*	8
CMT4 ¹	150	200	∞	0	1028.42*	12
CMT5 ¹	199	200	∞	0	1291.29*	16
CMT11 ¹	120	200	∞	0	1042.11*	7
CMT12 ¹	100	200	∞	0	819.56*	10
CMT6 ²	50	160	200	10	555.43	6
CMT7 ²	75	140	160	10	909.68	11
CMT8 ²	100	200	230	10	865.94	9
CMT9 ²	150	200	200	10	1162.55	14
CMT10 ²	199	200	200	10	1395.85	18
CMT13 ²	120	200	720	50	1541.14	11
CMT14 ²	100	200	1040	90	866.37	11

¹ The data of instances CMT1, CMT2, CMT3, CMT12, CMT11, CMT4 and CMT5 are the same of E-n51-k5, E-n76-k10, E-n101-k8, M-n101-k10, M-n121-k7, M-n151-k12 and M-n200-k16(k17), respectively. The only difference is the convention of not rounding the costs and not fixing the number of routes.

² Instances CMT6, CMT7, CMT8, CMT14, CMT13, CMT9 and CMT10 were generated by adding maximum route duration and service time to CMT1, CMT2, CMT3, CMT12, CMT11, CMT4 and CMT5, respectively. Vehicles are assumed to travel at unitary speed.

Table A.13

Instances of Golden et al. (1998).

Instance	n	Q	MD	BKS	K_{BKS}
G1 ⁴	240	550	650	5623.47	9
G2 ⁴	320	700	900	8404.61	10
G3 ⁴	400	900	1200	11036.22	10
G4 ⁴	480	1000	1600	13590.00	–
G5 ⁵	200	900	1800	6460.98	5
G6 ⁵	280	900	1500	8400.33	–
G7 ⁵	360	900	1300	10102.70	8
G8 ⁵	440	900	1200	11635.30	10
G9 ³	255	1000	∞	579.71	14
G10 ³	323	1000	∞	735.66	–
G11 ³	399	1000	∞	912.03	–
G12 ³	483	1000	∞	1101.50	–
G13 ²	252	1000	∞	857.19	26
G14 ²	320	1000	∞	1080.55*	30
G15 ²	396	1000	∞	1337.87	–
G16 ²	480	1000	∞	1611.56	–
G17 ¹	240	200	∞	707.76*	22
G18 ¹	300	200	∞	995.13*	27
G19 ¹	360	200	∞	1365.60*	33
G20 ¹	420	200	∞	1817.89	–

¹ Concentric six-pointed pointed stars.

² Concentric squares, depot in the center.

³ Concentric squares, depot in a corner.

⁴ Concentric circles.

⁵ Concentric rays.

in Table A.12. Column K_{BKS} gives the number of routes in the optimal/best known solution of each instance. Optimal solutions are marked with a *.

Table A.13 corresponds to the benchmark proposed in Golden et al. (1998). There are 12 CVRP instances and 8 instances for duration-constrained CVRP, the maximum durations (as there are no service times, this is equivalent to a bound on the maximum total distance traveled in a route) are given in column (MD). Table A.13 also gives the geometric patterns used for positioning the customers in each instance.

Table A.14 corresponds to a benchmark proposed in Rochat and Taillard (1995). Twelve instances from 75 to 150 customers were generated using a scheme where the depot is always in the center, the customers are clustered and demands are taken from an exponential distribution. An additional instance with 385 customers, obtained from real-world data, already appeared in Taillard (1993).

Table A.14

Instances of Rochat and Taillard (1995).

Instance	<i>n</i>	<i>Q</i>	BKS	<i>K</i> _{BKS}
tai75a ¹	75	1445	1618.36*	10
tai75b ¹	75	1679	1344.64*	9
tai75c ¹	75	1122	1291.01*	9
tai75d ¹	75	1699	1365.42*	9
tai100a ¹	100	1409	2041.34*	11
tai100b ¹	100	1842	1939.90*	11
tai100c ¹	100	2043	1406.20*	11
tai100d ¹	100	1297	1580.46*	11
tai150a ¹	150	1544	3055.23*	15
tai150b ¹	150	1918	2727.03*	14
tai150c ¹	150	2021	2358.66*	15
tai150d ¹	150	1874	2645.40*	14
tai385 ²	385	65	24366.41	47

¹ Customers are non-uniformly spread in several clusters. The number of clusters and their compactness are variable. Demands are generated following an exponential distribution.

² The data for the customers were generated based on real information from the canton of Vaud in Switzerland. Already appeared in Taillard (1993).

Table A.15

Instances of Li et al. (2005).

Instance	<i>n</i>	<i>Q</i>	MD	BKS	<i>K</i> _{BKS}
21	560	1200	1800	16212.74	–
22	600	900	1000	14499.04	15
23	640	1400	2200	18801.12	10
24	720	1500	2400	21389.33	–
25	760	900	900	16668.51	19
26	800	1700	2500	23971.74	–
27	840	900	900	17343.38	20
28	880	1800	2800	26565.92	–
29	960	2000	3000	29154.33	–
30	1040	2100	3200	31742.51	–
31	1120	2300	3500	34330.84	–
32	1200	2500	3600	37159.41	11

Note that a best value of 2341.84 for instance tai150c was mentioned on some benchmark instance repositories and then relayed in some papers. Yet, the exact algorithm of Pecin et al. (2014) determined that a solution of 2358.66 is optimal and most recent heuristics found the same value. We thus assume that this previous solution was erroneous.

Although there are no pure CVRP instances in this set, for the sake of completeness, Table A.15 presents a benchmark proposed in Li, Golden, and Wasil (2005), composed by 12 larger scale instances of the duration-constrained CVRP. Typical CVRP heuristics can be easily adapted to handle duration/distance constraints, and some papers on the CVRP also reported results in the Li benchmark. However, most recent state-of-the-art methods for the CVRP (Nagata & Bräysy, 2009; Vidal et al., 2012) did not. Therefore, in order to present solutions that correspond to the current point of algorithmic evolution, we performed runs with the algorithm of Vidal et al. (2012). Table A.15 has been updated to include the newly found best known solutions. We also highlight an issue related to the published solution of instance pr32, generated by a manual process in Li et al. (2005) with a distance value of 36919.24. This solution seems to have 10 routes from the figure in the paper, and as such cannot comply with the distance limit of 3600 units. For this reason, it was not included in the table.

Appendix B. Results on the instances of Li et al. (2005)

Table B.16 reports the results achieved with 10 runs of the hybrid genetic algorithm of Vidal et al. (2012), using the same parameter configuration as in the original paper and the same computing environment as other tests of this paper. The columns provide, in

Table B.16

Instances of Li et al. (2005) – results of Vidal et al. (2012).

Instance	<i>n</i>	Avg.	Best	<i>T</i> (min)
21	560	16212.83	16212.83	8.08
22	600	14545.76	14528.19	35.37
23	640	18826.22	18801.13	10.96
24	720	21389.43	21389.43	13.74
25	760	16753.56	16705.11	54.38
26	800	23977.73	23977.73	19.19
27	840	17411.96	17383.18	59.66
28	880	26566.04	26566.04	20.87
29	960	29154.34	29154.34	25.26
30	1040	31742.64	31742.64	30.11
31	1120	34330.94	34330.94	35.01
32	1200	37159.41	37159.41	51.07

turn, the instance number and size, the average and best solutions quality, and CPU time per run.

References

- Achuthan, N., Caccetta, L., & Hill, S. (2003). An improved branch-and-cut algorithm for the capacitated vehicle routing problem. *Transportation Science*, 37, 153–169.
- Araque, J., Kudva, G., Morin, T., & Pekny, J. (1994). A branch-and-cut algorithm for the vehicle routing problem. *Annals of Operations Research*, 50, 37–59.
- Augerat, P., Belenguer, J., Benavent, E., Corberán, A., Naddef, D., & Rinaldi, G. (1995). Computational results with a branch and cut code for the capacitated vehicle routing problem. *Technical report 949-M*. Grenoble, France: Université Joseph Fourier.
- Baldacci, R., Christofides, N., & Mingozzi, A. (2008). An exact algorithm for the vehicle routing problem based on the set partitioning formulation with additional cuts. *Mathematical Programming*, 115(2), 351–385.
- Baldacci, R., Mingozzi, A., & Roberti, R. (2011). New route relaxation and pricing strategies for the vehicle routing problem. *Operations Research*, 59(5), 1269–1283.
- Blasum, U., & Hochstättler, W. (2000). Application of the branch and cut method to the vehicle routing problem. *Technical report ZPR2000-386*. Zentrum für Angewandte Informatik Köln.
- Christofides, N., & Eilon, S. (1969). An algorithm for the vehicle-dispatching problem. *Operational Research Quarterly*, 20, 309–318.
- Christofides, N., Mingozzi, A., & Toth, P. (1979). The vehicle routing problem. In N. Christofides, A. Mingozzi, P. Toth, & C. Sandi (Eds.), *Combinatorial optimization: vol. 1* (pp. 315–338). Wiley Interscience.
- Clarke, G., & Wright, J. (1964). Scheduling of vehicles from a central depot to a number of delivery points. *Operations Research*, 12(4), 568–581.
- Contardo, C., & Martinelli, R. (2014). A new exact algorithm for the multi-depot vehicle routing problem under capacity and route length constraints. *Discrete Optimization*, 12, 129–146.
- Cook, W., Koch, T., Steffy, D. E., & Wolter, K. (2011). An exact rational mixed-integer programming solver. In *Integer programming and combinatorial optimization* (pp. 104–116). Springer.
- Dantzig, G. B., & Ramser, J. H. (1959). The truck dispatching problem. *Management Science*, 6(1), 80–91.
- Fisher, M. (1994). Optimal solution of vehicle routing problems using minimum K-trees. *Operations Research*, 42(4), 626–642.
- Frey, D., & Wang, H. (2006). Adaptive one-factor-at-a-time experimentation and expected value of improvement. *Technometrics*, 48(3), 418–431.
- Fukasawa, R., Longo, H., Lysgaard, J., Poggi de Aragão, M., Reis, M., Uchoa, E., et al. (2006). Robust branch-and-cut-and-price for the capacitated vehicle routing problem. *Mathematical Programming*, 106(3), 491–511.
- Gaskell, T. (1967). Bases for vehicle fleet scheduling. *Journal of the Operational Research Society*, 281–295.
- Gillett, B. E., & Miller, L. R. (1974). A heuristic algorithm for the vehicle-dispatch problem. *Operations Research*, 22(2), 340–349.
- Golden, B., Wasil, E., Kelly, J., & Chao, I. (1998). The impact of metaheuristics on solving the vehicle routing problem: Algorithms, problem sets, and computational results. In *Fleet management and logistics* (pp. 33–56). Springer.
- Gouveia, L. (1996). Multicommodity flow models for spanning trees with hop constraints. *European Journal of Operational Research*, 95(1), 178–190.
- Li, F., Golden, B., & Wasil, E. (2005). Very large-scale vehicle routing: New test problems, algorithms, and results. *Computers and Operations Research*, 32(5), 1165–1179.
- Lysgaard, J., Letchford, A., & Eglese, R. (2004). A new branch-and-cut algorithm for the capacitated vehicle routing problem. *Mathematical Programming*, 100, 423–445.
- Nagata, Y., & Bräysy, O. (2009). Edge assembly-based memetic algorithm for the capacitated vehicle routing problem. *Networks*, 54(4), 205–215.
- Pecin, D., Pessoa, A., Poggi, M., & Uchoa, E. (2014). Improved branch-cut-and-price for capacitated vehicle routing. In *Integer programming and combinatorial optimization* (pp. 393–403). Springer.

- Poggi, M., & Uchoa, E. (2014). New exact approaches for the capacitated VRP. In P. Toth, & D. Vigo (Eds.), *Vehicle routing: Problems, methods, and applications* (pp. 59–86). SIAM.
- Ralphs, T., Kopman, L., Pulleyblank, W., & Trotter, L. E. (2003). On the capacitated vehicle routing problem. *Mathematical Programming*, 94, 343–359.
- Rand, G. K. (2009). The life and times of the savings method for vehicle routing problems. *ORiON: The Journal of ORSSA*, 25(2), 125–145.
- Reinelt, G. (1991). TspLib – A traveling salesman problem library. *ORSA Journal on Computing*, 3(4), 376–384.
- Rochat, Y., & Taillard, É. D. (1995). Probabilistic diversification and intensification in local search for vehicle routing. *Journal of Heuristics*, 1(1), 147–167.
- Røpke, S. (2012). Branching decisions in branch-and-cut-and-price algorithms for vehicle routing problems. *Presentation in Column Generation*. <https://www.gerad.ca/colloques/ColumnGeneration2012/presentations/session7/Ropke.pdf>.
- Solomon, M. M. (1987). Algorithms for the vehicle routing and scheduling problems with time window constraints. *Operations Research*, 35(2), 254–265.
- Subramanian, A., Uchoa, E., & Ochi, L. S. (2013). A hybrid algorithm for a class of vehicle routing problems. *Computers and Operations Research*, 40(10), 2519–2531.
- Taillard, É. (1993). Parallel iterative search methods for vehicle routing problems. *Networks*, 23(8), 661–673.
- Vallada, E., Ruiz, R., & Framinan, J. M. (2015). New hard benchmark for flowshop scheduling problems minimising makespan. *European Journal of Operational Research*, 240(3), 666–677.
- Vidal, T., Crainic, T. G., Gendreau, M., Lahrichi, N., & Rei, W. (2012). A hybrid genetic algorithm for multidepot and periodic vehicle routing problems. *Operations Research*, 60(3), 611–624.
- Vidal, T., Crainic, T. G., Gendreau, M., & Prins, C. (2014). A unified solution framework for multi-attribute vehicle routing problems. *European Journal of Operational Research*, 234(3), 658–673.
- Wenger, K. (2003). *Generic cut generation methods for routing problems* Ph.D. thesis. Institute of Computer Science, University of Heidelberg.